

KEVIN XU

kxu9787@gmail.com | 530-601-3082 | San Francisco, CA | LinkedIn | GitHub

PROFESSIONAL SUMMARY

Built SQL, Python, and PySpark data pipelines and dashboarding at Deloitte for 3 years, supporting a DoD audit-readiness engagement covering a \$12B Air Force military equipment portfolio across 10+ programs. Translated technical financial requirements into reliable, repeatable calculation logic used in quarterly deliverables, and built reconciliation logic that traced a significant unaccounted variance in the team's financial data. M.A. in Statistics from UC Berkeley.

WORK EXPERIENCE

Risk & Financial Advisory Consultant

Deloitte, San Francisco | Promoted from Business Analyst, May 2024

Aug 2022 - Apr 2025

Supported a DoD Financial Improvement and Audit Readiness (FIAR) engagement from Apr 2023 to Apr 2025, building the data pipelines and calculation logic behind a \$12B Air Force military equipment portfolio.

- **Built and maintained canonical SQL-based ingestion pipelines** (Databricks, PySpark, Parquet) sourcing GAFS-BQ transaction data into an S3 data lake, plus an API-based tool automating input-template intake for 10+ Air Force programs (e.g., F-16); re-engineered the pipeline from ad hoc notebooks into modular, version-controlled components.
- **Translated complex military-equipment valuation requirements** (depreciation schedules by equipment type, lot-level percentage-complete tracking) into the calculation logic behind the team's quarterly Construction-in-Progress (CIP) deliverable, applying Military Equipment Valuation (MEV) methodology.
- **Wrote data-integrity checks reconciling** the team's automated deliverable against a legacy manual (Excel-based) process at the transaction level; authored the reconciliation logic, traced a significant unaccounted variance, and co-authored findings recommending priority programs for further review.
- **Owned technical documentation** (SOPs, schemas, data dictionaries) to standardize pipeline logic across the team, using Deloitte's internal ChatGPT-based LLM tool to accelerate drafting and debugging; mentored junior data scientists as the team scaled from 4 to 8 and the program expanded to Space Force.

From Sep 2022 to Apr 2023, supported a data-governance engagement for the U.S. Army on Deloitte's SEMOSS platform, building dashboards and analysis to help identify and sunset redundant Army policies.

- **Built graph-based clustering logic** organizing thousands of Army policies into a multi-level taxonomy (standards, clusters, systems) to identify redundant, sunset-able policies; visualized findings through interactive SEMOSS dashboards presented quarterly, giving Army leadership line of sight into policy redundancy.
- **Built additional dashboards** using SEMOSS, Python, SQL, and HTML, including an enterprise health grid tracking status and redundancy counts across the policy taxonomy's systems, hosted on AWS infrastructure.

Statistics Researcher

UC Davis Computational Communications Lab, Davis, CA

Aug - Nov 2021

- Applied canonical correlation and factor analysis across 31 psychological measures and 25 outcomes from 22 randomized experiments to identify traits predicting resilience against algorithmic manipulation; contributed to peer-reviewed research under PI Martin Hilbert.

PROJECTS

Feature Cost vs. Signal in Credit Risk Prediction

Jan - May 2026

- Built and stress-tested 11 model specifications and regularization paths (LASSO/Ridge) across 30K accounts to run scenario and sensitivity analysis on which cost drivers justified their data-acquisition cost; a 2-variable model matched full-model accuracy, and regime-dependent testing showed cost drivers shift by borrower lifecycle stage; together, these findings informed a trade-off recommendation to cut scope without sacrificing signal.

Forecasting: CA Industrial Emissions

Sep - Dec 2025

- Forecasted a trend-dominated cost series using SARIMA, Lasso, and Gradient Boosting; SARIMA had the strongest residual diagnostics (Ljung-Box $p > 0.75$), while a simplified 4-feature alternative preserved comparable accuracy.

EDUCATION

University of California, Berkeley

Master of Arts, Statistics | Elizabeth Scott Excellence Scholarship (merit)

Sep 2024 - May 2026

University of California, Davis

Bachelor of Science, Statistics | Dean's List

Sep 2020 - Jun 2022

SKILLS

Languages: SQL (PostgreSQL, T-SQL), Python (pandas, NumPy, scikit-learn, XGBoost, LightGBM), PySpark, R, Bash

Data Engineering & Tools: ETL/ELT Pipelines, Distributed Computing (Spark), S3 Data Lake, Databricks, MLflow, Tableau, Power BI, Matplotlib, Seaborn, AWS (S3, Glue, Athena, SageMaker), Git

Financial & Analytical: Financial Calculations, Variance Analysis, Scenario Analysis, Forecasting (SARIMA), Regression (OLS/WLS/GLS), Causal Inference