

Kevin Xu

kxu9787@gmail.com | 530-601-3082 | San Francisco, CA | [LinkedIn](#) | [GitHub](#) | [Website](#)

Professional Summary

Built and re-engineered SQL and PySpark ETL pipelines at Deloitte for 3 years, integrating structured and semi-structured data into an S3 data lake for a \$12B Air Force audit-remediation program. Traced a significant unaccounted variance through reconciliation logic I wrote. M.A. in Statistics, UC Berkeley.

Work Experience

Risk & Financial Advisory Consultant

Deloitte, San Francisco | Promoted from Business Analyst, May 2024

Aug 2022 - Apr 2025

Supported a DoD Financial Improvement and Audit Remediation (FIAR) engagement from Apr 2023 to Apr 2025, building the data pipelines and calculation logic behind a \$12B Air Force military equipment portfolio.

- **Built and maintained two ingestion pipelines feeding an S3 data lake:** a SQL-based pipeline (AWS Glue, Databricks, PySpark, Parquet) for structured GAFS-BQ transaction data, and an API-based tool automating intake of progress reports across 10+ programs (e.g., F-35 aircraft). Re-engineered both pipelines from ad hoc notebooks into modular, version-controlled components.
- **Built Tableau and Power BI dashboards** from pipeline outputs for quarterly client deliverables, internal pipeline-health monitoring, and program-level compliance tracking.
- **Modeled military-equipment valuation logic** for the team's Construction-in-Progress (CIP) deliverable, updating calculation logic as underlying business rules evolved.
- **Wrote data-integrity checks reconciling** the team's automated deliverable against a legacy manual (Excel-based) process at the transaction level; authored the reconciliation logic and traced a significant unaccounted variance.
- **Owned technical documentation** (SOPs, schemas, data dictionaries) to standardize pipeline logic across the team, using Deloitte's internal ChatGPT-based LLM tool to accelerate drafting and debugging; mentored junior data scientists as the team scaled from 4 to 8 and the program expanded to Space Force.

From Sep 2022 to Apr 2023, supported a data-governance engagement for the U.S. Army on Deloitte's SEMOSS platform, building dashboards and analysis to help identify and sunset redundant Army policies.

- **Built graph-based clustering logic** organizing thousands of Army policies into a multi-level taxonomy (standards, clusters, systems) to identify redundant, sunset-able policies; visualized findings through interactive SEMOSS dashboards presented quarterly, giving Army leadership line of sight into policy redundancy.
- **Built additional dashboards** using SEMOSS, Python, SQL, and HTML, including an enterprise health grid tracking status and redundancy counts across the policy taxonomy's systems, hosted on AWS infrastructure.

Statistics Researcher

UC Davis Computational Communications Lab, Davis, CA

Aug - Nov 2021

- Applied canonical correlation and factor analysis across 31 psychological measures and 25 outcomes from 22 randomized experiments to identify traits predicting resilience against algorithmic manipulation; contributed to peer-reviewed research under PI Martin Hilbert.

Projects

Feature Cost vs. Signal in Credit Risk Prediction

Jan - May 2026

- Conducted cost-efficiency tradeoff analysis across 11 regularized model specifications (LASSO/Ridge) on 30k credit card accounts to identify which features justified their data acquisition cost.
- Regime-dependent testing showed cost drivers shift by borrower lifecycle stage. Findings informed a trade-off recommendation to cut scope without sacrificing signal.

Forecasting: CA Industrial Emissions

Sep - Dec 2025

- Compared SARIMA, Lasso and Gradient Boosting to identify the most robust forecasting approach for a trend-dominated cost series.
- SARIMA had the strongest residual diagnostics (Ljung-Box $p > 0.75$), while a simplified 4-feature alternative preserved comparable accuracy.

Education

University of California, Berkeley

Sep 2024 - May 2026

Master of Arts, Statistics | Elizabeth Scott Excellence Scholarship (merit)

University of California, Davis

Sep 2020 - Jun 2022

Bachelor of Science, Statistics | Dean's List

Skills

Languages: SQL (PostgreSQL, T-SQL), Python (pandas, NumPy, scikit-learn, XGBoost, LightGBM), PySpark, R, Bash

Data Engineering & Tools: ETL/ELT Pipelines, Data Modeling (dbt), Distributed Computing (Spark), S3 Data Lake, Databricks, MLflow, Tableau, Power BI, Matplotlib, Seaborn, AWS (S3, Glue, Athena, SageMaker), Git

Financial & Analytical: Financial Calculations, Variance Analysis, Scenario Analysis, Forecasting (SARIMA)